

training

January 3, 2025

```
[1]: import numpy as np
import pandas as pd
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df = pd.read_csv('/kaggle/input/credit-card/creditcard.csv')
df.head()
```

```
[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	

	V8	V9	...	V21	V22	V23	V24	V25	\
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

```
[3]: df.shape
```

```
[3]: (284807, 31)
```

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
```

Data columns (total 31 columns):

#	Column	Non-Null Count	Dtype
0	Time	284807 non-null	float64
1	V1	284807 non-null	float64
2	V2	284807 non-null	float64
3	V3	284807 non-null	float64
4	V4	284807 non-null	float64
5	V5	284807 non-null	float64
6	V6	284807 non-null	float64
7	V7	284807 non-null	float64
8	V8	284807 non-null	float64
9	V9	284807 non-null	float64
10	V10	284807 non-null	float64
11	V11	284807 non-null	float64
12	V12	284807 non-null	float64
13	V13	284807 non-null	float64
14	V14	284807 non-null	float64
15	V15	284807 non-null	float64
16	V16	284807 non-null	float64
17	V17	284807 non-null	float64
18	V18	284807 non-null	float64
19	V19	284807 non-null	float64
20	V20	284807 non-null	float64
21	V21	284807 non-null	float64
22	V22	284807 non-null	float64
23	V23	284807 non-null	float64
24	V24	284807 non-null	float64
25	V25	284807 non-null	float64
26	V26	284807 non-null	float64
27	V27	284807 non-null	float64
28	V28	284807 non-null	float64
29	Amount	284807 non-null	float64
30	Class	284807 non-null	int64

dtypes: float64(30), int64(1)

memory usage: 67.4 MB

```
[5]: df.isnull().sum()
```

```
[5]: Time      0
     V1        0
     V2        0
     V3        0
     V4        0
     V5        0
     V6        0
     V7        0
```

```
V8          0
V9          0
V10         0
V11         0
V12         0
V13         0
V14         0
V15         0
V16         0
V17         0
V18         0
V19         0
V20         0
V21         0
V22         0
V23         0
V24         0
V25         0
V26         0
V27         0
V28         0
Amount      0
Class       0
dtype: int64
```

```
[6]: df['Class'].value_counts()
```

```
[6]: Class
0    284315
1      492
Name: count, dtype: int64
```

```
[7]: import matplotlib.pyplot as plt
import seaborn as sns

# Visualize class distribution
sns.countplot(x='Class', data=df)
plt.title('Class Distribution (Fraud vs Non-Fraud)')
plt.xlabel('Class')
plt.ylabel('Count')
plt.show()
```



```
[8]: legit_df = df[df['Class'] == 0]
      fraud_df = df[df['Class'] == 1]

      legit_df.shape
```

```
[8]: (284315, 31)
```

```
[9]: fraud_df.shape
```

```
[9]: (492, 31)
```

```
[10]: legit_df.Amount.describe()
```

```
[10]: count    284315.000000
      mean      88.291022
      std     250.105092
      min       0.000000
      25%       5.650000
      50%      22.000000
      75%      77.050000
```

```
max      25691.160000
Name: Amount, dtype: float64
```

```
[11]: fraud_df.Amount.describe()
```


```
[11]: count      492.000000
      mean      122.211321
      std       256.683288
      min        0.000000
      25%        1.000000
      50%        9.250000
      75%       105.890000
      max      2125.870000
      Name: Amount, dtype: float64
```

```
[12]: df.groupby('Class').mean().T
```


```
[12]: Class          0          1
Time    94838.202258  80746.806911
V1        0.008258   -4.771948
V2       -0.006271    3.623778
V3        0.012171   -7.033281
V4       -0.007860    4.542029
V5        0.005453   -3.151225
V6        0.002419   -1.397737
V7        0.009637   -5.568731
V8       -0.000987    0.570636
V9        0.004467   -2.581123
V10       0.009824   -5.676883
V11      -0.006576    3.800173
V12       0.010832   -6.259393
V13       0.000189   -0.109334
V14       0.012064   -6.971723
V15       0.000161   -0.092929
V16       0.007164   -4.139946
V17       0.011535   -6.665836
V18       0.003887   -2.246308
V19      -0.001178    0.680659
V20      -0.000644    0.372319
V21      -0.001235    0.713588
V22      -0.000024    0.014049
V23       0.000070   -0.040308
V24       0.000182   -0.105130
V25      -0.000072    0.041449
V26      -0.000089    0.051648
V27      -0.000295    0.170575
V28      -0.000131    0.075667
```

Amount 88.291022 122.211321


```
[13]: # Distribution of 'Amount' feature
sns.displot(df['Amount'], bins=25, kde=True, color='blue')
plt.title('Transaction Amount Distribution')
plt.xlabel('Amount')
plt.ylabel('Frequency')
plt.show()
```




```
[14]: plt.figure(figsize=(8,5))
ax = sns.distplot(legit_df['Amount'],label='fraudulent',hist=False)
ax = sns.distplot(fraud_df['Amount'],label='non fraudulent',hist=False)
ax.set(xlabel='Seconds elapsed between the transection and the first transection')
plt.show()
```



```
[15]: # Distribution of 'Time' feature
sns.displot(df['Time'], bins=50, kde=True, color='green')
plt.title('Transaction Time Distribution')
plt.xlabel('Time')
plt.ylabel('Frequency')
plt.show()
```




```
[16]: plt.figure(figsize=(8,5))
      ax = sns.distplot(legit_df['Time'],label='fraudulent',hist=False)
      ax = sns.distplot(fraud_df['Time'],label='non fraudulent',hist=False)
      ax.set(xlabel='Seconds elapsed between the transection and the first transection')
      plt.show()
```

```
[17]: df.drop('Time', axis=1, inplace=True)
```

```
[18]: plt.figure(figsize=(20, 12))
sns.heatmap(df.corr(), cmap='coolwarm', annot=True, fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()
```

```
[20]: # Calculate outliers for each numeric feature
numeric_cols = df.select_dtypes(include=['float64', 'int64']).columns
numeric_cols = numeric_cols.drop('Class')

outliers_by_feature = {}
for col in numeric_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1

    outliers = ((df[col] < (Q1 - 1.5 * IQR)) | (df[col] > (Q3 + 1.5 * IQR)))
    outliers_by_feature[col] = outliers.sum()

print(f'Outliers in {col}: {outliers.sum()}')
```

```
Outliers in V1: 7062
Outliers in V2: 13526
Outliers in V3: 3363
Outliers in V4: 11148
Outliers in V5: 12295
Outliers in V6: 22965
Outliers in V7: 8948
Outliers in V8: 24134
Outliers in V9: 8283
Outliers in V10: 9496
```

```


Outliers in V11: 780
Outliers in V12: 15348
Outliers in V13: 3368
Outliers in V14: 14149
Outliers in V15: 2894
Outliers in V16: 8184
Outliers in V17: 7420
Outliers in V18: 7533
Outliers in V19: 10205
Outliers in V20: 27770
Outliers in V21: 14497
Outliers in V22: 1317
Outliers in V23: 18541
Outliers in V24: 4774
Outliers in V25: 5367
Outliers in V26: 5596
Outliers in V27: 39163
Outliers in V28: 30342
Outliers in Amount: 31904

```

```

[21]: # Visualize outliers
plt.figure(figsize=(12,6))
plt.bar(outliers_by_feature.keys(), outliers_by_feature.values())
plt.xticks(rotation=45)
plt.title('Number of Outliers by Feature')
plt.ylabel('Number of Outliers')
plt.tight_layout()
plt.show()

```



```
[22]: from sklearn.preprocessing import RobustScaler
```

```
scaler = RobustScaler()
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
df.head()
```

```
[22]:
```

	V1	V2	V3	V4	V5	V6	V7	\
0	-0.616237	-0.098602	1.228905	0.878152	-0.217859	0.631245	0.177406	
1	0.524929	0.143100	-0.006970	0.293974	0.087726	0.164395	-0.105740	
2	-0.615587	-1.002407	0.830932	0.251024	-0.344345	1.778007	0.668164	
3	-0.440239	-0.178789	0.841250	-0.529808	0.033775	1.303832	0.175637	
4	-0.526089	0.579239	0.713861	0.265632	-0.270695	0.317183	0.491625	

	V8	V9	V10	...	V21	V22	V23	V24	\
0	0.142432	0.334787	0.185689	...	0.026866	0.253109	-0.320791	0.032681	
1	0.117064	-0.164482	-0.074854	...	-0.473332	-0.602719	0.363442	-0.479557	
2	0.420388	-1.179796	0.303796	...	0.668917	0.714254	2.974603	-0.919589	
3	0.662489	-1.076888	0.038374	...	-0.190105	-0.001408	-0.578786	-1.531963	
4	-0.546463	0.700808	0.855099	...	0.048266	0.739092	-0.407980	0.126293	

	V25	V26	V27	V28	Amount	Class
0	0.167619	-0.241182	0.816731	-0.246091	1.783274	0
1	0.225462	0.313475	-0.063781	0.026519	-0.269825	0
2	-0.515430	-0.153111	-0.350218	-0.540962	4.983721	0
3	0.944482	-0.298959	0.379163	0.382611	1.418291	0
4	-0.333308	0.976221	1.347133	1.553716	0.670579	0

[5 rows x 30 columns]

```
[23]: # Random State
random_state = 1
```

```
[24]: X = df.drop('Class', axis=1)
y = df['Class']
```

```
[25]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=random_state, stratify=y)

# Split Data Shape
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[25]: ((227845, 29), (56962, 29), (227845,), (56962,))
```

```
[26]: print("\nClass Distribution Before Resampling:")
print(pd.Series(y_train).value_counts())
```

Class Distribution Before Resampling:

Class

0 227451

1 394

Name: count, dtype: int64

```
[27]: from imblearn.over_sampling import SMOTE

smote = SMOTE(random_state=random_state)
X_resampled, y_resampled = smote.fit_resample(X_train, y_train)

print("\nClass Distribution After Resampling:")
print(pd.Series(y_resampled).value_counts())
```

Class Distribution After Resampling:

Class

0 227451

1 227451

Name: count, dtype: int64

```
[28]: from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import LinearSVC
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier
from sklearn.calibration import CalibratedClassifierCV
```

```
[29]: from sklearn.metrics import *
import pickle

all_metrics = {}

def evaluate_model(model, X_train, y_train, X_test, y_test, model_name,
    ↪threshold=0.99):
    model.fit(X_train, y_train)
    y_pred_proba = model.predict_proba(X_test)[: , 1]
    y_pred = (y_pred_proba >= threshold).astype(int)

    # Calculate metrics
    accuracy = accuracy_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_pred_proba)

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred)
```

```

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title(f'Confusion Matrix - {model_name}')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.show()

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'ROC Curve - {model_name}')
plt.legend()
plt.show()

# Precision-Recall Curve
precision, recall, thresholds = precision_recall_curve(y_test, y_pred_proba)
pr_auc = auc(recall, precision)
plt.figure(figsize=(8, 6))
plt.plot(recall, precision, label=f'PR Curve (AUC = {pr_auc:.4f})')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title(f'Precision-Recall Curve - {model_name}')
plt.legend()
plt.show()

# Optimal threshold (e.g., maximizing F1-Score)
f1_scores = 2 * (precision * recall) / (precision + recall + 1e-10)
optimal_idx = np.argmax(f1_scores)
optimal_threshold = thresholds[optimal_idx]
print(f"Optimal Threshold (Max F1-Score): {optimal_threshold:.4f}")

metrics = {
    'Accuracy': accuracy,
    'ROC-AUC': roc_auc,
    'Precision': precision_score(y_test, y_pred),
    'Recall': recall_score(y_test, y_pred),
    'F1-Score': f1_score(y_test, y_pred),
    'Optimal Threshold': optimal_threshold,
    'True Negatives': cm[0][0],
    'False Positives': cm[0][1],
    'False Negatives': cm[1][0],
    'True Positives': cm[1][1]
}

```

```

all_metrics[model_name] = metrics

print(f"\n{model_name} Classification Report:")
print(classification_report(y_test, y_pred))
metrics_df = pd.DataFrame([metrics], index=[model_name])

# Save the model
model_path = f'{model_name.lower().replace(" ", "_")} _model.pkl'
with open(model_path, 'wb') as file:
    pickle.dump(model, file)
print(f"Model saved at: {model_path}")


return metrics_df.T

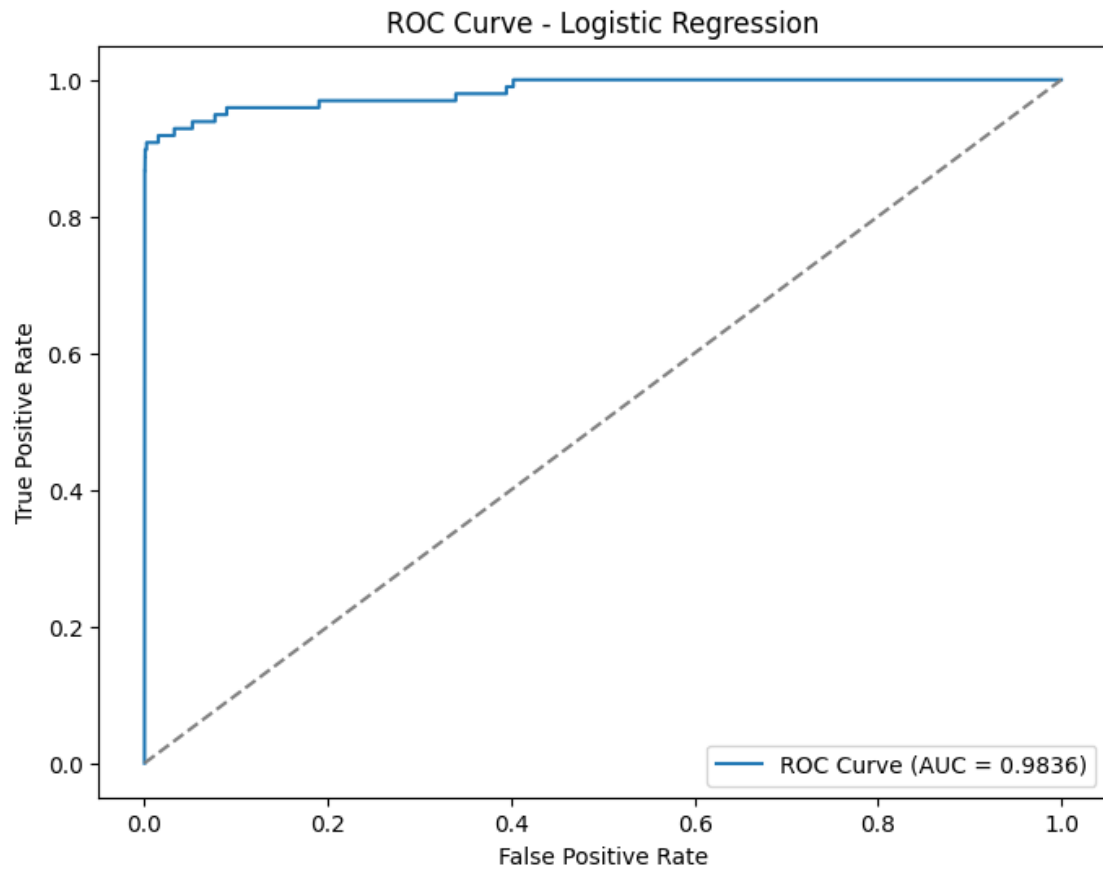
```

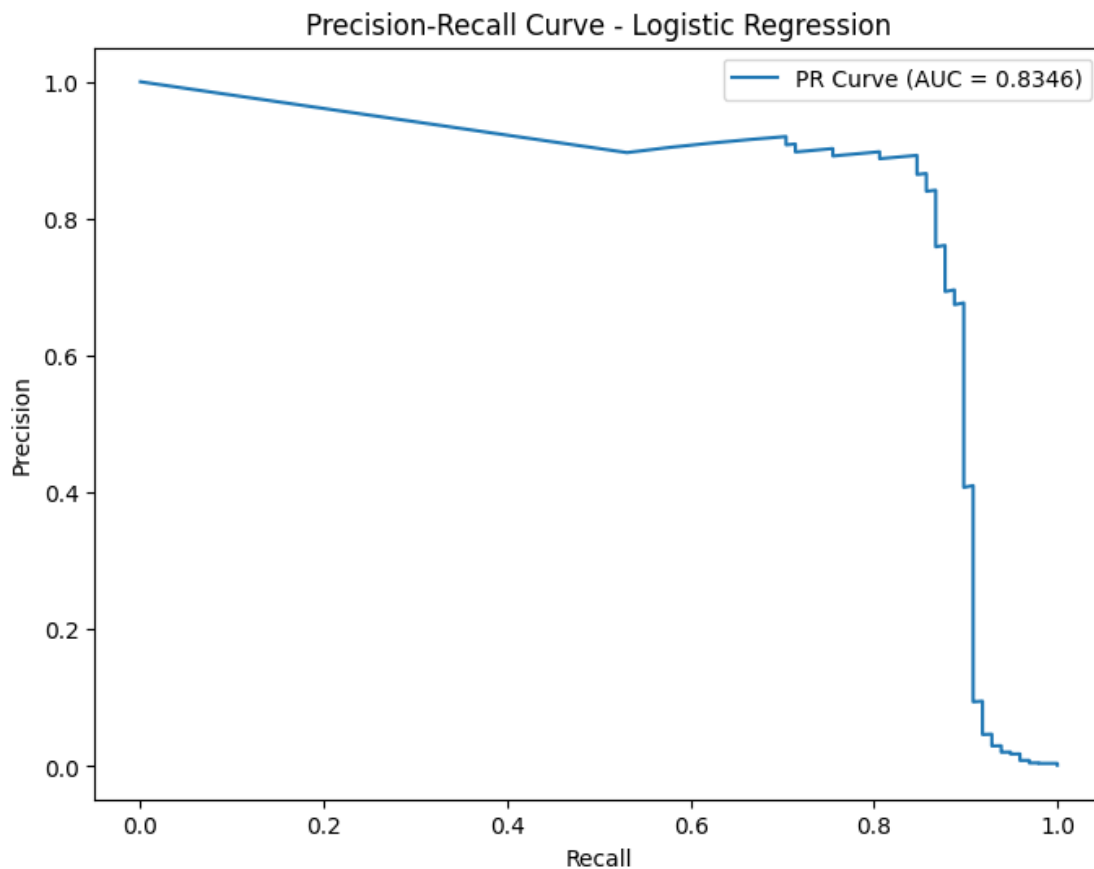
```

[30]: # Logistic Regression
logistic_model = LogisticRegression(random_state=random_state)
evaluate_model(logistic_model, X_resampled, y_resampled, X_test, y_test,
↳ "Logistic Regression")

```







Optimal Threshold (Max F1-Score): 1.0000

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.62	0.90	0.73	98
accuracy			1.00	56962
macro avg	0.81	0.95	0.86	56962
weighted avg	1.00	1.00	1.00	56962


Model saved at: logistic_regression_model.pkl


[30]:


	Logistic Regression
Accuracy	0.998859
ROC-AUC	0.983626
Precision	0.615385
Recall	0.897959
F1-Score	0.730290

Optimal Threshold	1.000000
True Negatives	56809.000000
False Positives	55.000000
False Negatives	10.000000
True Positives	88.000000

```
[31]: # Decision Tree
dt_model = DecisionTreeClassifier(random_state=random_state)
evaluate_model(dt_model, X_resampled, y_resampled, X_test, y_test, "Decision_
↪Tree")
```







Optimal Threshold (Max F1-Score): 1.0000

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.41	0.89	0.56	98
accuracy			1.00	56962
macro avg	0.70	0.94	0.78	56962
weighted avg	1.00	1.00	1.00	56962

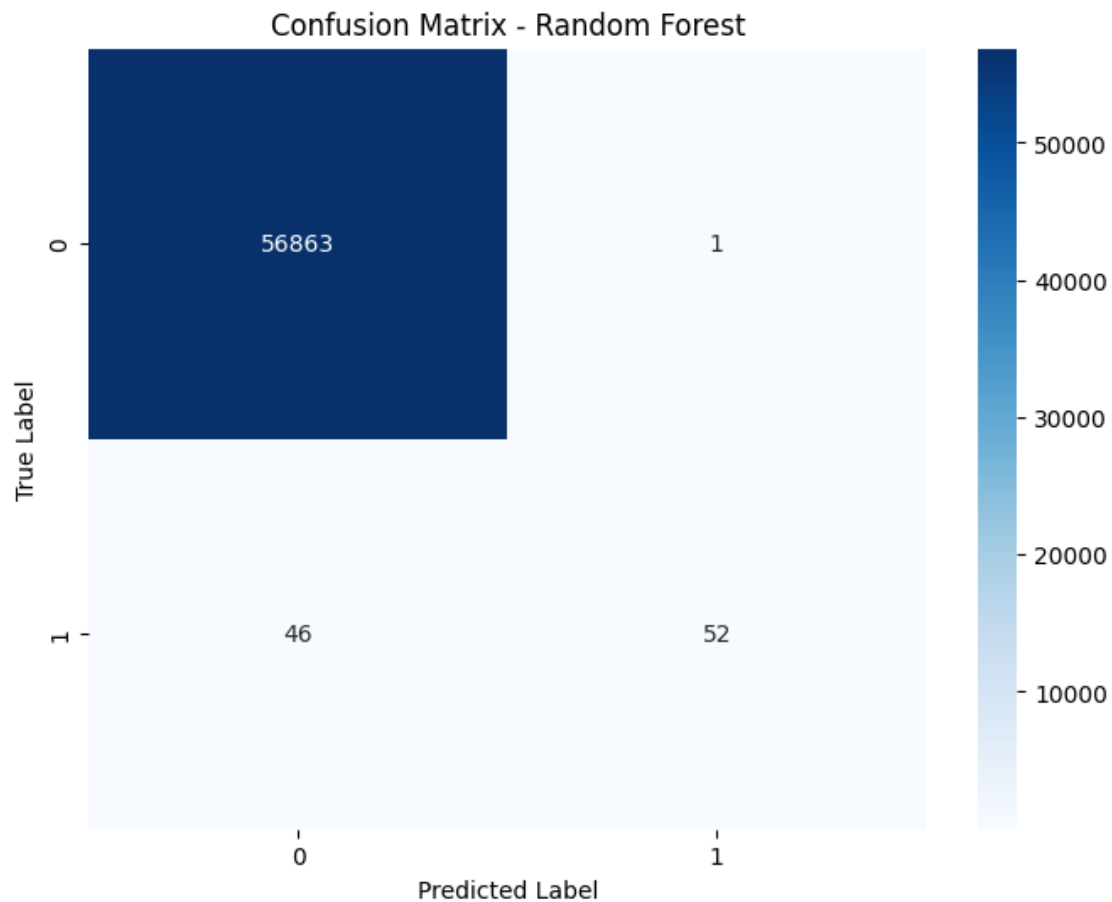
Model saved at: decision_tree_model.pkl


[31]:

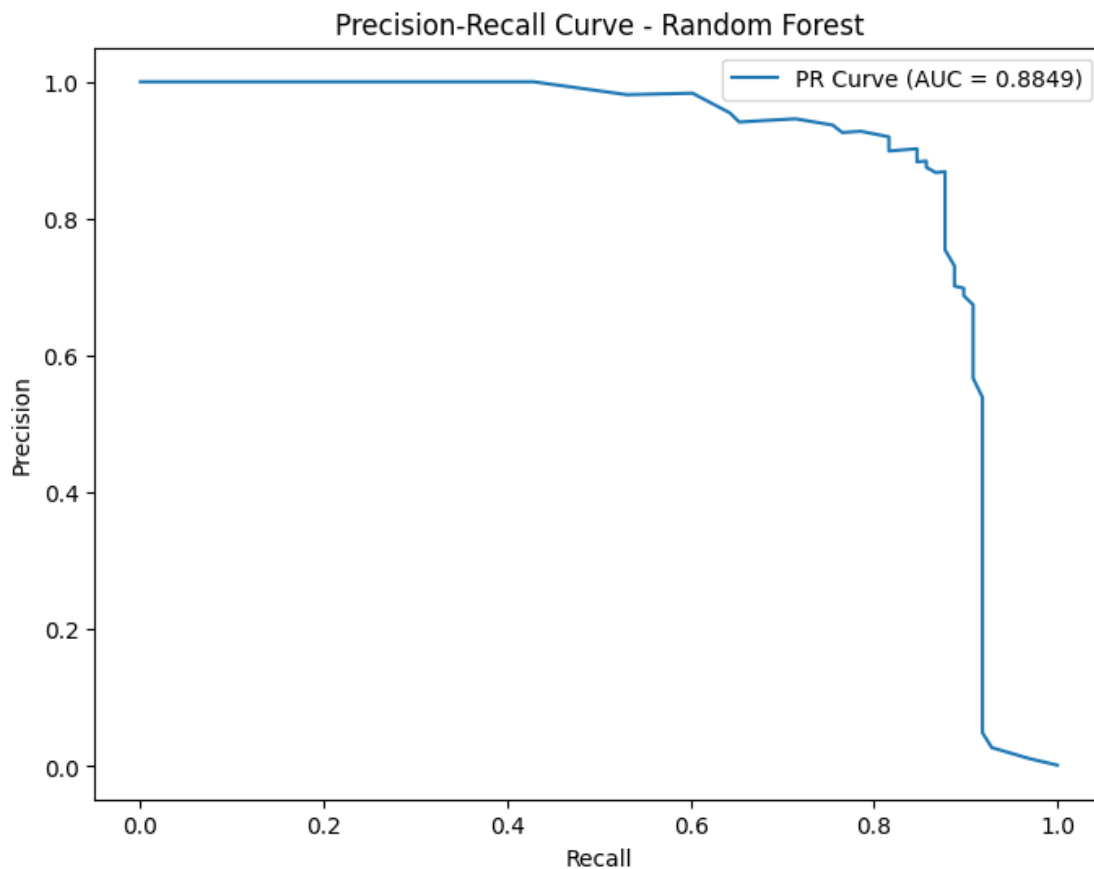
	Decision Tree
Accuracy	0.997577
ROC-AUC	0.942761
Precision	0.406542
Recall	0.887755
F1-Score	0.557692

```
Optimal Threshold      1.000000
True Negatives         56737.000000
False Positives        127.000000
False Negatives        11.000000
True Positives         87.000000
```

```
[32]: # Random Forest
rf_model = RandomForestClassifier(random_state=random_state)
evaluate_model(rf_model, X_resampled, y_resampled, X_test, y_test, "Random_
↪Forest")
```







Optimal Threshold (Max F1-Score): 0.7500

Random Forest Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.98	0.53	0.69	98
accuracy			1.00	56962
macro avg	0.99	0.77	0.84	56962
weighted avg	1.00	1.00	1.00	56962


Model saved at: random_forest_model.pkl


[32]:

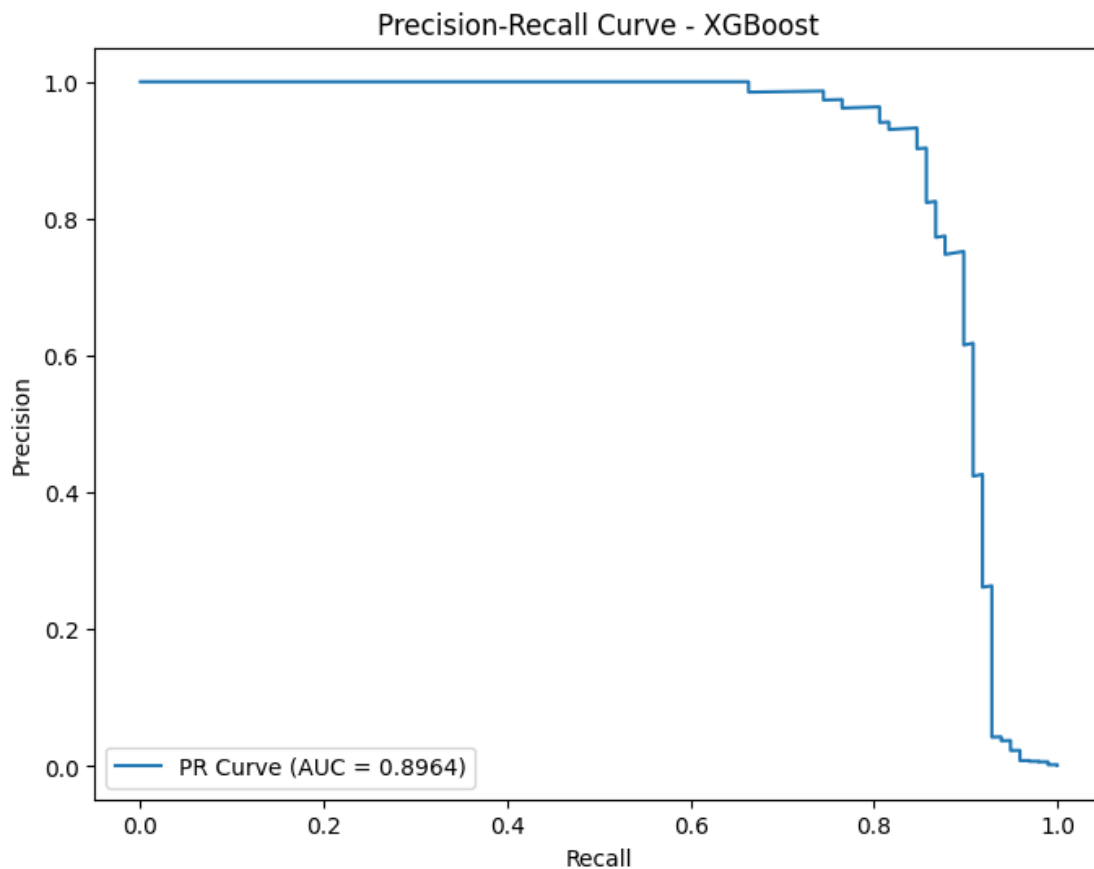
	Random Forest
Accuracy	0.999175
ROC-AUC	0.977802
Precision	0.981132
Recall	0.530612
F1-Score	0.688742


```
Optimal Threshold      0.750000
True Negatives         56863.000000
False Positives        1.000000
False Negatives        46.000000
True Positives         52.000000
```

```
[33]: # XGBoost
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
    random_state=random_state)
evaluate_model(xgb_model, X_resampled, y_resampled, X_test, y_test, "XGBoost")
```







Optimal Threshold (Max F1-Score): 0.9976

XGBoost Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.91	0.85	0.88	98
accuracy			1.00	56962
macro avg	0.96	0.92	0.94	56962
weighted avg	1.00	1.00	1.00	56962

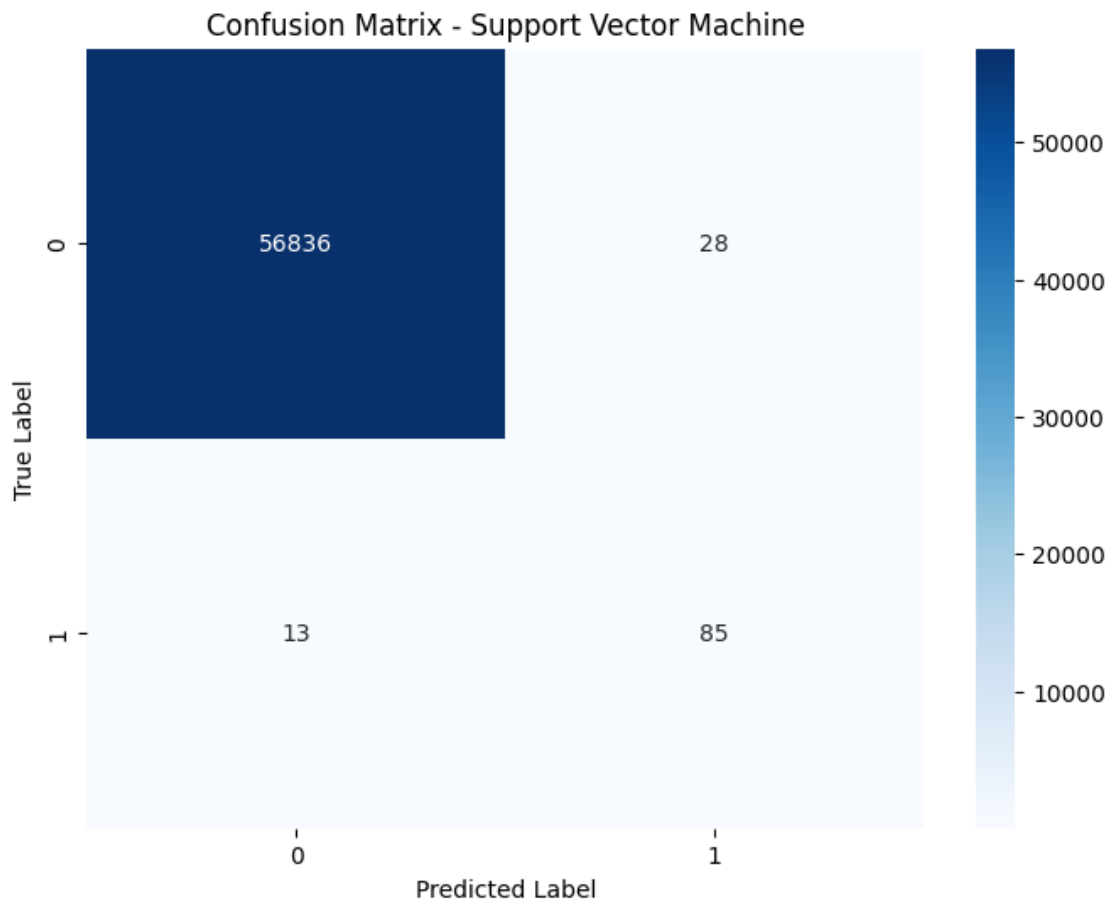
Model saved at: xgboost_model.pkl

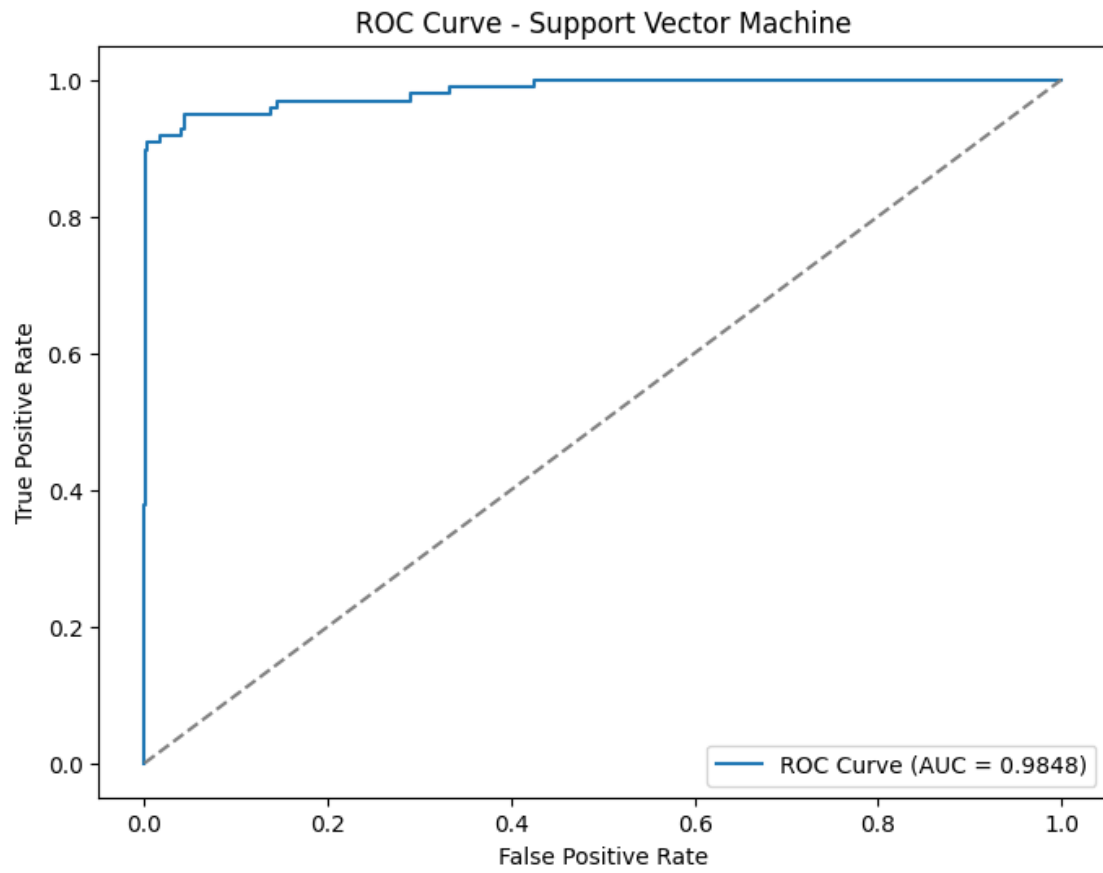
[33]:

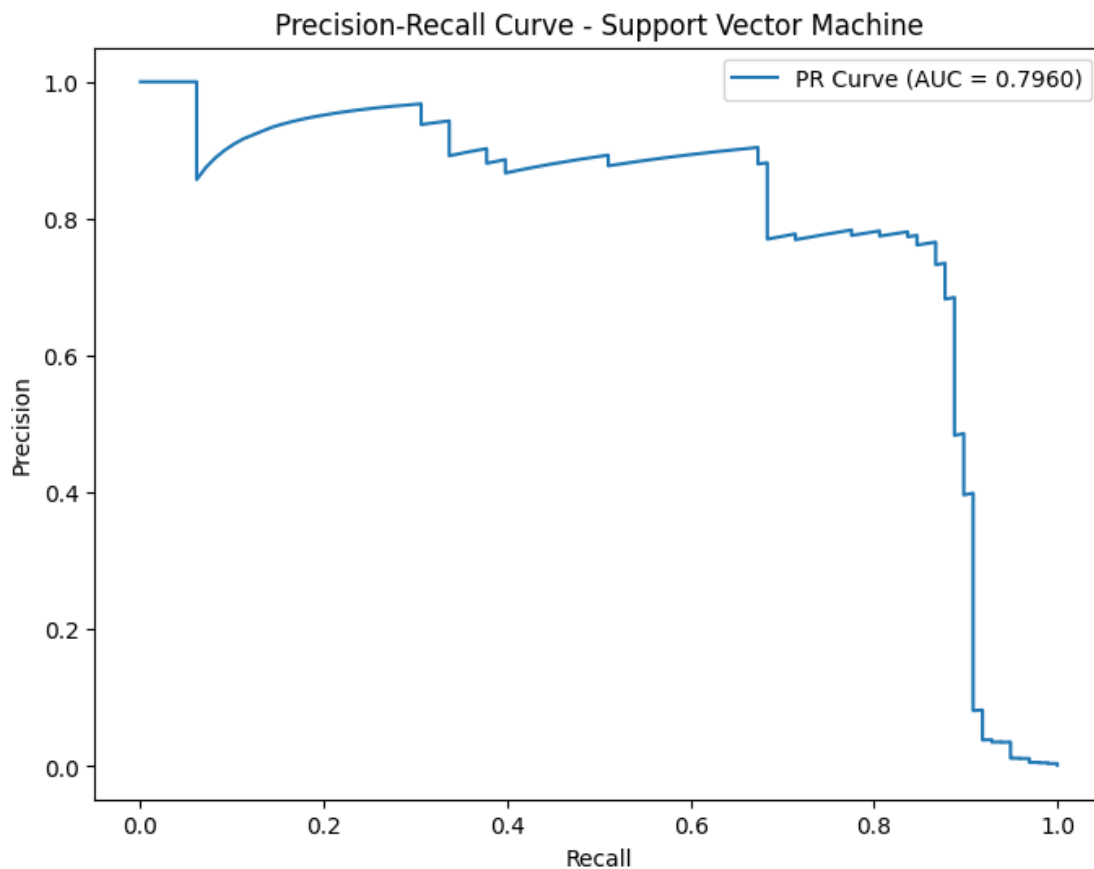
	XGBoost
Accuracy	0.999596
ROC-AUC	0.984043
Precision	0.912088
Recall	0.846939
F1-Score	0.878307

```
Optimal Threshold      0.997589
True Negatives        56856.000000
False Positives        8.000000
False Negatives        15.000000
True Positives         83.000000
```

```
[34]: # Support Vector Machine
linear_svc = LinearSVC(max_iter=10000)
svm_model = CalibratedClassifierCV(base_estimator=linear_svc, cv=5)
evaluate_model(svm_model, X_resampled, y_resampled, X_test, y_test, "Support_
↳Vector Machine")
```







Optimal Threshold (Max F1-Score): 0.9945

Support Vector Machine Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.75	0.87	0.81	98
accuracy			1.00	56962
macro avg	0.88	0.93	0.90	56962
weighted avg	1.00	1.00	1.00	56962

Model saved at: support_vector_machine_model.pkl


[34]: Support Vector Machine

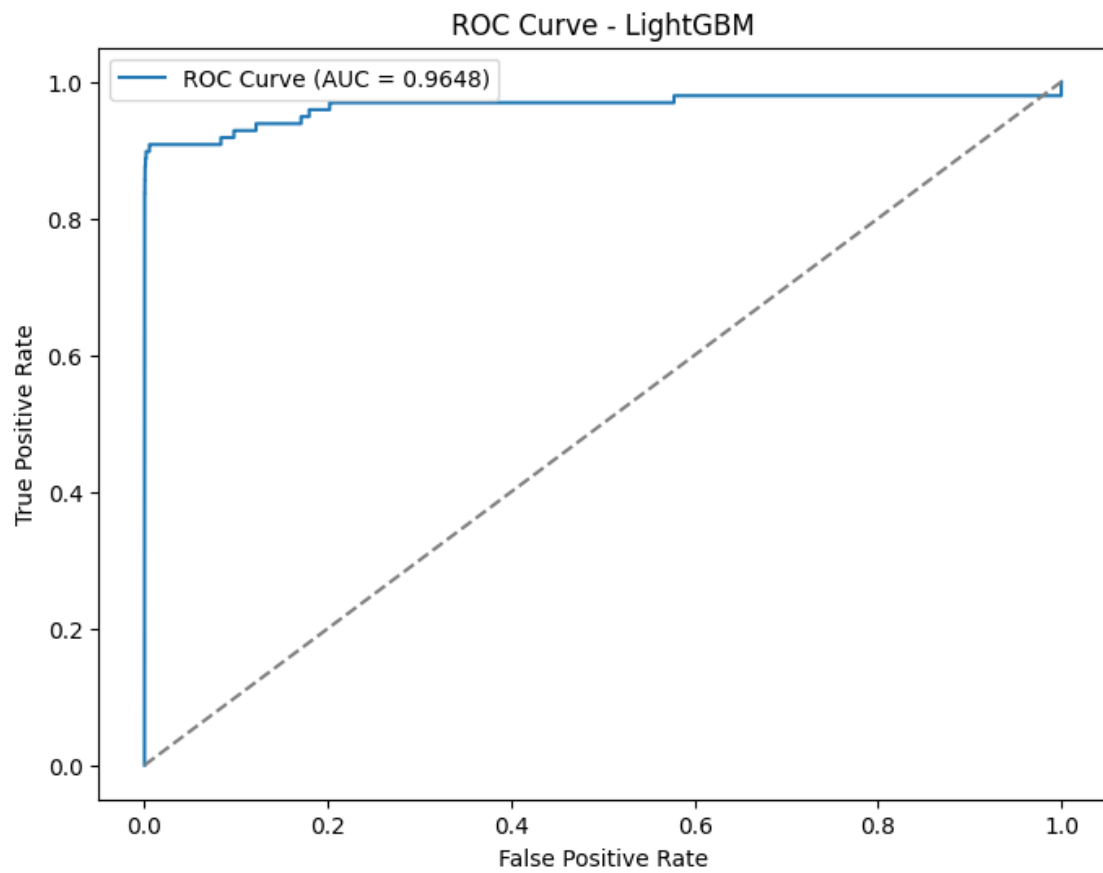
Accuracy	0.999280
ROC-AUC	0.984772
Precision	0.752212
Recall	0.867347
F1-Score	0.805687

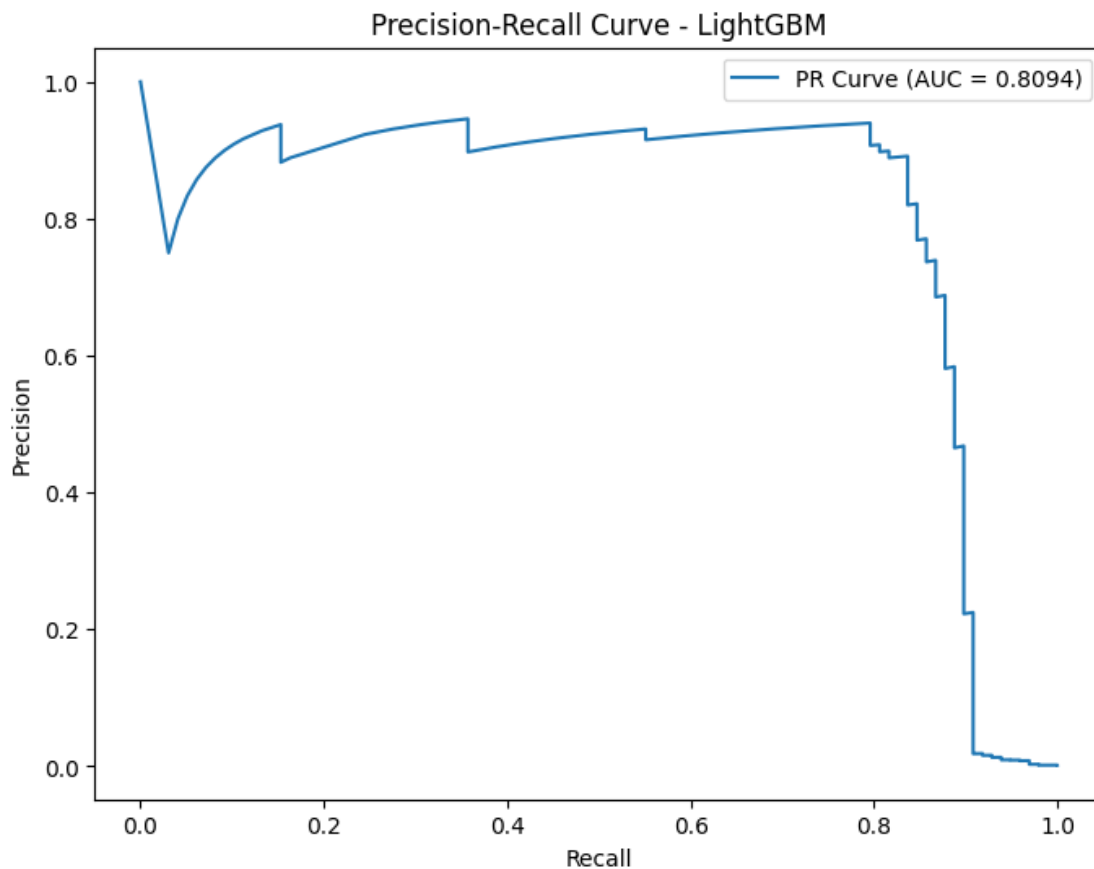
Optimal Threshold	0.994470
True Negatives	56836.000000
False Positives	28.000000
False Negatives	13.000000
True Positives	85.000000

```
[35]: # LightGBM
lgbm_model = LGBMClassifier(random_state=random_state)
evaluate_model(lgbm_model, X_resampled, y_resampled, X_test, y_test, "LightGBM")
```

[LightGBM] [Info] Number of positive: 227451, number of negative: 227451
 [LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.076268 seconds.
 You can set `force_col_wise=true` to remove the overhead.
 [LightGBM] [Info] Total Bins 7395
 [LightGBM] [Info] Number of data points in the train set: 454902, number of used features: 29
 [LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000







Optimal Threshold (Max F1-Score): 0.9935

LightGBM Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.88	0.84	0.86	98
accuracy			1.00	56962
macro avg	0.94	0.92	0.93	56962
weighted avg	1.00	1.00	1.00	56962

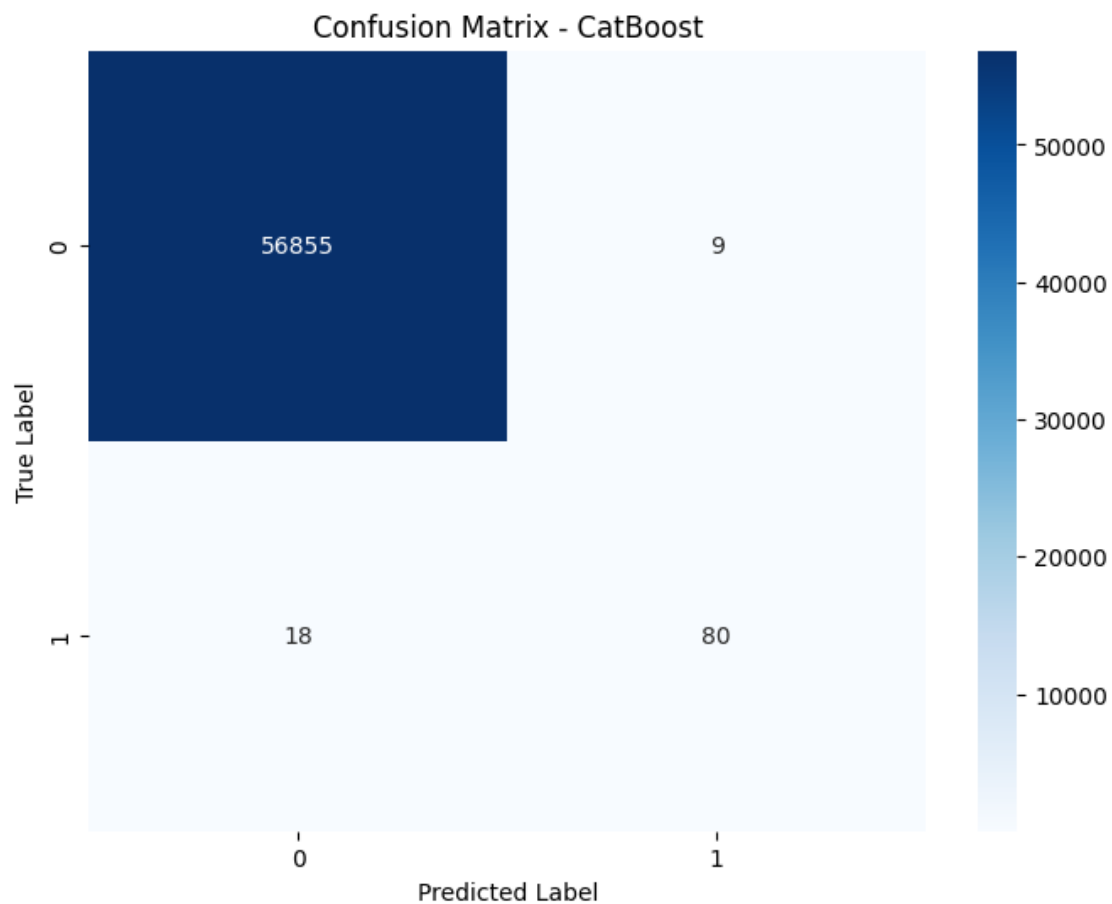
Model saved at: lightgbm_model.pkl


[35]:


	LightGBM
Accuracy	0.999526
ROC-AUC	0.964819
Precision	0.881720
Recall	0.836735
F1-Score	0.858639

Optimal Threshold 0.993473
True Negatives 56853.000000
False Positives 11.000000
False Negatives 16.000000
True Positives 82.000000

```
[36]: # CatBoost  
catboost_model = CatBoostClassifier(verbose=0, random_state=random_state)  
evaluate_model(catboost_model, X_resampled, y_resampled, X_test, y_test, ↪"CatBoost")
```







Optimal Threshold (Max F1-Score): 0.9572

CatBoost Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.90	0.82	0.86	98
accuracy			1.00	56962
macro avg	0.95	0.91	0.93	56962
weighted avg	1.00	1.00	1.00	56962

Model saved at: catboost_model.pkl


[36]:

	CatBoost
Accuracy	0.999526
ROC-AUC	0.981731
Precision	0.898876
Recall	0.816327
F1-Score	0.855615

Optimal Threshold 0.957154
True Negatives 56855.000000
False Positives 9.000000
False Negatives 18.000000
True Positives 80.000000

```
[37]: def compare_models(metrics_list):  
      # Convert metrics to DataFrame  
      metrics_df = pd.DataFrame(metrics_list)  
  
      # Plot metrics comparison  
      metrics_to_plot = ['Accuracy', 'ROC-AUC', 'Precision', 'Recall', 'F1-Score']  
      metrics_df.T[metrics_to_plot].plot(kind='bar', figsize=(12, 8))  
      plt.title('Model Performance Comparison')  
      plt.ylabel('Score')  
      plt.xticks(rotation=45)  
      plt.legend(loc='lower right')  
      plt.tight_layout()  
      plt.show()  
  
      return metrics_df
```

```
[38]: result_df = compare_models(all_metrics)  
result_df
```



```
[38]:
```

	Logistic Regression	Decision Tree	Random Forest	\
Accuracy	0.998859	0.997577	0.999175	
ROC-AUC	0.983626	0.942761	0.977802	
Precision	0.615385	0.406542	0.981132	
Recall	0.897959	0.887755	0.530612	
F1-Score	0.730290	0.557692	0.688742	
Optimal Threshold	1.000000	1.000000	0.750000	
True Negatives	56809.000000	56737.000000	56863.000000	
False Positives	55.000000	127.000000	1.000000	
False Negatives	10.000000	11.000000	46.000000	
True Positives	88.000000	87.000000	52.000000	

	XGBoost	Support Vector Machine	LightGBM	\
Accuracy	0.999596	0.999280	0.999526	
ROC-AUC	0.984043	0.984772	0.964819	
Precision	0.912088	0.752212	0.881720	
Recall	0.846939	0.867347	0.836735	
F1-Score	0.878307	0.805687	0.858639	
Optimal Threshold	0.997589	0.994470	0.993473	
True Negatives	56856.000000	56836.000000	56853.000000	
False Positives	8.000000	28.000000	11.000000	
False Negatives	15.000000	13.000000	16.000000	
True Positives	83.000000	85.000000	82.000000	

	CatBoost
Accuracy	0.999526
ROC-AUC	0.981731
Precision	0.898876
Recall	0.816327
F1-Score	0.855615
Optimal Threshold	0.957154
True Negatives	56855.000000
False Positives	9.000000
False Negatives	18.000000
True Positives	80.000000

```
[39]: # Find the best model based on Accuracy
metrics_df = result_df.T
best_model = metrics_df.loc[metrics_df['Accuracy'].idxmax()]
print("\nBest Model:")
print(best_model)
```

```
Best Model:
Accuracy      0.999596
ROC-AUC       0.984043
```

Precision	0.912088
Recall	0.846939
F1-Score	0.878307
Optimal Threshold	0.997589
True Negatives	56856.000000
False Positives	8.000000
False Negatives	15.000000
True Positives	83.000000

Name: XGBoost, dtype: float64